

Lab 5: RV64 缺页异常处理

3210105488 刘扬

1 实验过程

1.1 准备工作

从repo同步框架，不赘述。

在 `trap_handler` 中传入新的参数 `stval` 供本次实验使用：

```
void trap_handler(unsigned long scause, unsigned long sepc, unsigned long stval, struct pt_regs *regs);
```

1.2 实现VMA

1.2.1 修改 `proc.h`

结构体定义如下：

```
#define VM_X_MASK          0x0000000000000008
#define VM_W_MASK          0x0000000000000004
#define VM_R_MASK          0x0000000000000002
#define VM_ANONYM          0x0000000000000001

struct vm_area_struct {
    uint64 vm_start;          /* VMA 对应的用户态虚拟地址的开始 */
    uint64 vm_end;            /* VMA 对应的用户态虚拟地址的结束 */
    uint64 vm_flags;          /* VMA 对应的 flags */

    /* uint64_t file_offset_on_disk */ /* 原本需要记录对应的文件在磁盘上的位置，
                                       但是我们只有一个文件 uapp，所以暂时不需要记录 */

    uint64 vm_content_offset_in_file; /* 如果对应了一个文件，
                                       那么这块 VMA 起始地址对应的文件内容相对文件起始位置的偏移量，
                                       也就是 ELF 中各段的 p_offset 值 */

    uint64 vm_content_size_in_file; /* 对应的文件内容的长度。
                                       思考为什么还需要这个域？
                                       和 (vm_end-vm_start)
                                       一比，不是冗余了吗？ */
};

/* 线程数据结构 */
struct task_struct {
    struct thread_info thread_info;
    uint64 state; // 线程状态
    uint64 counter; // 运行剩余时间
```

```

uint64 priority; // 运行优先级 1最低 10最高
uint64 pid;      // 线程id

struct thread_struct thread;
uint64 satp;
pagetable_t pgd;
uint64 vma_cnt; /* 下面这个数组里的元素的数量 */
struct vm_area_struct vmas[0]; /* 为什么可以开大小为 0 的数组?
                                这个定义可以和前面的 vma_cnt 换个位置吗? */
};

```

添加 `do_mmap` 和 `vm_area_struct` 声明:

```

void do_mmap(struct task_struct *task, uint64_t addr, uint64_t length, uint64_t flags,
             uint64_t vm_content_offset_in_file, uint64_t vm_content_size_in_file);

struct vm_area_struct *find_vma(struct task_struct *task, uint64_t addr);

```

1.2.2 修改 `task_init`

修改 `task_init` 函数代码, 更改为 Demand Paging:

- 取消之前实验中对 `U-MODE` 代码以及栈进行的映射
- 调用 `do_mmap` 函数, 建立用户 `task` 的虚拟地址空间信息。

在下述代码中, `load_program` 调用了 `do_mmap` 建立用户 `task` 的虚拟地址空间信息。

```

void task_init() {
    test_init(NR_TASKS);
    ...
    for (int i = 1; i < NR_TASKS; i++) {
        task[i] = (struct task_struct*)kalloc();
        task[i]->state = TASK_RUNNING;
        ...
        load_program(task[i], pgtbl); // 在load_program中调用do_mmap
        uint64 satp = (uint64)0x8 << 60;
        satp |= ((uint64)(pgtbl) - PA2VA_OFFSET) >> 12;
        task[i]->satp = satp;
    }
    printk("...proc_init done!\n");
}

```

`load_program` 如下:

```

static uint64 load_program(struct task_struct* task, pagetable_t pgtbl) {
    Elf64_Ehdr* ehdr = (Elf64_Ehdr*)_sramdisk;

    uint64 phdr_start = (uint64)ehdr + ehdr->e_phoff;
    int phdr_cnt = ehdr->e_phnum;
}

```

```

Elf64_Phdr* phdr;
for (int i = 0; i < phdr_cnt; i++) {
    phdr = (Elf64_Phdr*)(phdr_start + sizeof(Elf64_Phdr) * i);
    if (phdr->p_type == PT_LOAD) {
        uint64 offset = (uint64)(phdr->p_vaddr) % PGSIZE;
        uint64 num_pages = (phdr->p_memsz + offset) / PGSIZE;
        if (num_pages * PGSIZE < (phdr->p_memsz + offset)){
            num_pages++;
        }
        uint64 length = num_pages * PGSIZE;
        uint64 flags = phdr->p_flags << 1;
        do_mmap(task, phdr->p_vaddr, length, flags, phdr->p_offset, phdr->p_filesz);
    }
}
do_mmap(task, USER_END-PGSIZE, PGSIZE, VM_R_MASK | VM_W_MASK | VM_ANONYM, 0, 0);
task->thread.sepc = ehdr->e_entry;
unsigned long sstatus = csr_read(sstatus);
sstatus &= ~(1 << 8); // sstatus[SPP] = 0
sstatus |= 1 << 5; // sstatus[SPIE] = 1
sstatus |= 1 << 18; // sstatus[SUM] = 1
task->thread.sstatus = sstatus;
task->thread.sscratch = USER_END;
}

```

在 `load_program` 中只建立用户 `task` 的虚拟地址空间信息，并不进行读取，当真正访问该页时，才读取对应的页。

1.2.3 do_mmap 实现

简单地将传入的参数进行存储即可。

```

void do_mmap(struct task_struct *task, uint64 addr, uint64 length, uint64 flags,
            uint64 vm_content_offset_in_file, uint64 vm_content_size_in_file) {
    struct vm_area_struct* vma = &(task->vmas[task->vma_cnt]);
    task->vma_cnt++;
    vma->vm_flags = flags;
    vma->vm_content_offset_in_file = vm_content_offset_in_file;
    vma->vm_content_size_in_file = vm_content_size_in_file;
    vma->vm_start = addr;
    vma->vm_end = addr + length;
}

```

在实现完上述步骤后，`trap_handler` 会检测到 page fault，但还不能对 page fault 进行处理。

1.2.4 实现 Page Fault 的检测与处理

在 `trap_handler` 中检测 page fault

```

void trap_handler(unsigned long scause, unsigned long sepc, unsigned long stval, struct
pt_regs *regs) {

```

```

    if (scause == 0x8000000000000005) {
        ...
    }
    else if (scause == 8) {
        ...
    }
    else if (scause == (uint64)0xc || scause == (uint64)0xd || scause == (uint64)0xf) {
        printk("[S] Supervisor Page Fault, scause: %lx, stval: %lx, sepc: %lx\n", scause,
stval, sepc);
        do_page_fault(regs, stval);
    }
    else {
        ...
    }
}

```

调用 `do_page_fault` 处理page fault:

- 需要先找到错误地址(stval)的vma, 调用了 `find_vma`

```

struct vm_area_struct *find_vma(struct task_struct *task, uint64_t addr){
    for(int i = 0; i < task->vma_cnt; i++){
        if(addr >= task->vmas[i].vm_start && addr < task->vmas[i].vm_end){
            return &(task->vmas[i]);
        }
    }
    return NULL;
}

```

- 调用 `create_mapping` 进行页表映射

```

create_mapping(current->pgd, PGROUNDDOWN(stval), (uint64)new_page-PA2VA_OFFSET,
PGSIZE, (vma->vm_flags | 0x11))

```

- 如果是非匿名的page, 则需要将 `uapp` 对应地址的内容拷贝进来。注意要做对齐, 如果是vm_start的所在的第一页, 做拷贝时需要把vm_start的页偏移量考虑上, 如果是后面的页, 则拷贝一整页。

```

if(!(vma->vm_flags & VM_ANONYM)){
    char *src_addr = (char*)(_sramdisk + vma->vm_content_offset_in_file);
    if(stval - PGROUNDDOWN(vma->vm_start) < PGSIZE){
        uint64 offset = stval % PGSIZE;
        for(int j = 0; j < PGSIZE - offset; j++){
            new_page[j + offset] = src_addr[j];
        }
    }
    else{
        uint64 pg_offset = (stval - PGROUNDDOWN(vma->vm_start)) / PGSIZE;
        uint64 offset = stval % PGSIZE;
        for(int j = 0; j < PGSIZE; j++){
            new_page[j] = src_addr[j + PGSIZE * pg_offset - offset];
        }
    }
}

```

```

    }
}
}

```

do_page_fault 完整代码如下:

```

void do_page_fault(struct pt_regs *regs, unsigned long stval) {
    struct vm_area_struct *vma = find_vma(current, stval);
    char *new_page = alloc_page();
    create_mapping(current->pgd, PGROUNDDOWN(stval), (uint64)new_page-PA2VA_OFFSET,
PGSIZE, (vma->vm_flags | 0x11));
    if(!(vma->vm_flags & VM_ANONYM)){
        char *src_addr = (char*)(_sramdisk + vma->vm_content_offset_in_file);
        if(stval - PGROUNDDOWN(vma->vm_start) < PGSIZE){
            uint64 offset = stval % PGSIZE;
            for(int j = 0;j < PGSIZE - offset;j++){
                new_page[j + offset] = src_addr[j];
            }
        }
        else{
            uint64 pg_offset = (stval - PGROUNDDOWN(vma->vm_start)) / PGSIZE;
            uint64 offset = stval % PGSIZE;
            for(int j = 0;j < PGSIZE;j++){
                new_page[j] = src_addr[j + PGSIZE * pg_offset - offset];
            }
        }
    }
}
}

```

2 实验结果

我在实验中创建了4个user进程，每个进程发生3次page fault，依次分别是取指令错误，栈错误和访存错误，一共12次page fault。

```
switch to [PID = 1 PRIORITY = 37 COUNTER = 4]
[S] Supervisor Page Fault, scause: 0000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010124
[S] Supervisor Page Fault, scause: 0000000000000000d, stval: 0000000000011880, sepc: 0
0000000000010140
[PID = 1] is running, variable: 0
[PID = 1] is running, variable: 1
[PID = 1] is running, variable: 2
[PID = 1] is running, variable: 3
switch to [PID = 4 PRIORITY = 66 COUNTER = 5]
[S] Supervisor Page Fault, scause: 0000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010124
[S] Supervisor Page Fault, scause: 0000000000000000d, stval: 0000000000011880, sepc: 0
0000000000010140
[PID = 4] is running, variable: 0
[PID = 4] is running, variable: 1
[PID = 4] is running, variable: 2
[PID = 4] is running, variable: 3
[PID = 4] is running, variable: 4
[PID = 4] is running, variable: 5
```

```
switch to [PID = 3 PRIORITY = 52 COUNTER = 8]
[S] Supervisor Page Fault, scause: 0000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010124
[S] Supervisor Page Fault, scause: 0000000000000000d, stval: 0000000000011880, sepc: 0
0000000000010140
[PID = 3] is running, variable: 0
[PID = 3] is running, variable: 1
[PID = 3] is running, variable: 2
[PID = 3] is running, variable: 3
[PID = 3] is running, variable: 4
[PID = 3] is running, variable: 5
[PID = 3] is running, variable: 6
[PID = 3] is running, variable: 7
[PID = 3] is running, variable: 8
switch to [PID = 2 PRIORITY = 88 COUNTER = 9]
[S] Supervisor Page Fault, scause: 0000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010124
[S] Supervisor Page Fault, scause: 0000000000000000d, stval: 0000000000011880, sepc: 0
0000000000010140
[PID = 2] is running, variable: 0
[PID = 2] is running, variable: 1
[PID = 2] is running, variable: 2
[PID = 2] is running, variable: 3
```

```
[PID = 2] is running, variable: 4
[PID = 2] is running, variable: 5
[PID = 2] is running, variable: 6
[PID = 2] is running, variable: 7
[PID = 2] is running, variable: 8
[PID = 2] is running, variable: 9
SET [PID = 4 PRIORITY = 66 COUNTER = 1]
SET [PID = 3 PRIORITY = 52 COUNTER = 4]
SET [PID = 2 PRIORITY = 88 COUNTER = 10]
SET [PID = 1 PRIORITY = 37 COUNTER = 4]
switch to [PID = 4 PRIORITY = 66 COUNTER = 1]
[PID = 4] is running, variable: 6
switch to [PID = 1 PRIORITY = 37 COUNTER = 4]
[PID = 1] is running, variable: 4
[PID = 1] is running, variable: 5
[PID = 1] is running, variable: 6
[PID = 1] is running, variable: 7
[PID = 1] is running, variable: 8
switch to [PID = 3 PRIORITY = 52 COUNTER = 4]
[PID = 3] is running, variable: 9
```

```
switch to [PID = 1 PRIORITY = 37 COUNTER = 4]
[PID = 1] is running, variable: 4
[PID = 1] is running, variable: 5
[PID = 1] is running, variable: 6
[PID = 1] is running, variable: 7
[PID = 1] is running, variable: 8
switch to [PID = 3 PRIORITY = 52 COUNTER = 4]
[PID = 3] is running, variable: 9
[PID = 3] is running, variable: 10
[PID = 3] is running, variable: 11
[PID = 3] is running, variable: 12
switch to [PID = 2 PRIORITY = 88 COUNTER = 10]
[PID = 2] is running, variable: 10
[PID = 2] is running, variable: 11
[PID = 2] is running, variable: 12
[PID = 2] is running, variable: 13
[PID = 2] is running, variable: 14
[PID = 2] is running, variable: 15
QEMU: Terminated
```

```
parallels@ubuntu-linux-22-04-02-desktop:~/oslab/lab5$
```

3 思考题

1. `uint64_t vm_content_size_in_file;` 对应的文件内容的长度。为什么还需要这个域?

在 `do_mapp` 中, `vm_content_size_in_file` 映射的是 `uapp` 中的 `phdr->p_filesz`, 而本段内容在内容镜像中的大小是大于 `phdr->p_filesz` 的。在上个lab的思考题中也提到过这二者的区别: 段的文件映像是指存储在文件中的段的部分。段的内存映像是指在程序执行期间加载到内存中的段的部分。`p_memsz` (即 `length`) 大于 `p_filesz` (即 `vm_content_size_in_file`) 的原因是, 可加载段可能包含 `.bss` 节, 该节包含未初始化的数据。将此数据存储在磁盘上会很浪费, 因此, 仅在ELF文件加载到内存后才占用空间。`struct`

2. `vm_area_struct vmas[0];` 为什么可以开大小为 0 的数组? 这个定义可以和前面的 `vma_cnt` 换个位置吗?

参考: [你了解C语言的“柔性数组”吗? 看完你就知道了 - 知乎 \(zhihu.com\)](https://www.zhihu.com/question/20462034)

该数组是柔性数组。这个特性允许你在定义结构体的时候创建一个空数组, 而这个数组的大小可以在程序运行的过程中根据你的需求进行更改。特别注意的一点是: 这个空数组必须声明为结构体的最后一个成员, 并且还要求这样的结构体至少包含一个其他类型的成员。

可以看到柔性数组必须在结构体的最后声明。在本实验中, 如果和 `vma_cnt` 交换位置, 在后续写入 `vmas` 数组时, 会将 `vma_cnt` 的值覆盖掉, 导致错误。