

Lab 6: 实现 fork 机制

3210105488 刘扬

1 实验过程

1.1 准备工作

基于lab5，不赘述。

因为这次实验只初始化一个线程，需要修改 `task_init()`

```
task[1] = (struct task_struct*)kalloc();
task[1]->state = TASK_RUNNING;
task[1]->counter = 0;
task[1]->priority = 1;
task[1]->pid = 1;

task[1]->thread.ra = (uint64)__dummy;
task[1]->thread.sp = (uint64)(task[1]) + PGSIZE;
task[1]->thread_info.kernel_sp = (unsigned long)(task[1])+PGSIZE;

pagetable_t pgtbl = (pagetable_t)kalloc();
for(int j = 0; j < PGSIZE / sizeof(uint64*); j++){
    pgtbl[j] = swapper_pg_dir[j];
}
task[1]->pgd = pgtbl;

load_program(task[1], pgtbl);
uint64 satp = (uint64)0x8 << 60;
satp |= ((uint64)(pgtbl) - PA2VA_OFFSET) >> 12;
task[1]->satp = satp;
printk("[S] Initialized: pid: 1, priority: 1, counter: 0\n");
```

同时要注意在 `schedule()` 时，因为有些线程没有初始化，需要额外判断 `task[i] != NULL`

1.2 实现 `sys_clone()`

1.2.1 参考 `task_init` 创建一个新的 task，将的 parent task 的整个页复制到新创建的 `task_struct` 页上。将 `thread.ra` 设置为 `__ret_from_fork`，并正确设置 `thread.sp`

复制task_struct并设置thread.ra

```
struct task_struct *child = (struct task_struct*)kalloc();
for(int i = 0; i < PGSIZE; i++){
    ((char*)child)[i] = ((char*)current)[i];
}
child->thread.ra = (uint64)__ret_from_fork;
```

设置thread.sp: 这里的sp是用于给 `__switch_to` 正确的栈指针, 使 `__ret_from_fork` 的栈指针是正确的。

```
uint64 offset = (uint64)regs % PGSIZE;
struct pt_regs *child_regs = (struct pt_regs*)(child + offset);
//__switch_to->__ret_from_fork(in _traps)->user program
child->thread.sp = (uint64)child_regs; // for __switch_to
```

1.2.2 利用参数 regs 来计算出 child task 的对应的 pt_regs 的地址, 并将其中的 a0, sp, sepc 设置成正确的值

通过parent task的regs相对于页起始地址的偏移量算出child task的pt_regs的地址

```
uint64 offset = (uint64)regs % PGSIZE;
struct pt_regs *child_regs = (struct pt_regs*)(child + offset);

for(int i = 0; i < 37 * 8; i++){
    ((char*)child_regs)[i] = ((char*)regs)[i];
}
child_regs->x[2] = (uint64)child_regs; // for _traps
child_regs->x[10] = 0; //return 0 (child thread)
child_regs->sepc = regs->sepc + 4;
```

1.2.3 为 child task 申请 user stack, 并将 parent task 的 user stack数据复制到其中

```
uint64 child_stack = alloc_page();
for(int i = 0; i < PGSIZE; i++){
    ((char*)child_stack)[i] = ((char*)(USER_END - PGSIZE))[i];
}
```

1.2.4 同时将子 task 的 user stack 的地址保存在 thread_info->user_sp 中

```
child->thread_info.kernel_sp = (uint64)(child) + PGSIZE;
child->thread_info.user_sp = (uint64)USER_END;
```

1.2.5 为 child task 分配一个根页表, 并仿照 setup_vm_final 来创建内核空间的映射

```
child->pgd = (uint64*)alloc_page();
uint64 satp = (uint64)0x8 << 60;
satp |= ((uint64)(child->pgd) - PA2VA_OFFSET) >> 12;
child->satp = satp;
for(int i = 0; i < PGSIZE; i++){
    ((char*)child->pgd)[i] = ((char*)swapper_pg_dir)[i];
}
create_mapping(child->pgd, USER_END - PGSIZE, (uint64)child_stack - PA2VA_OFFSET, PGSIZE, 23);
```

1.2.6 根据 parent task 的页表和 vma 来分配并拷贝 child task 在用户态会用到的内存

遍历parent task的vma，对于每个vma中的每一页，通过走页表来判断parent task是否有做映射，如果做了映射，child task的相应地址也需要做映射

```
void copy_vma(uint64* pgd){
    for (int i = 0; i < current->vma_cnt; i++) {
        struct vm_area_struct *cur_vma = &(current->vmas[i]);
        uint64 cur_addr = cur_vma->vm_start;
        while (cur_addr < cur_vma->vm_end) {
            walk_pgd(current->pgd, PGROUNDDOWN(cur_addr), pgd);
            cur_addr += PGSIZE;
        }
    }
}
```

walk_pgd的主要逻辑是判断parent task的每个page table entry的V位，如果V位都被置1，说明该页已经映射，相应的child task需要做拷贝

```
int walk_pgd(uint64* parent_pgd, uint64 va, uint64* child_pgd){
    uint64 cur_vpn;
    uint64* cur_pgtbl;
    uint64 cur_pte;
    cur_pgtbl = parent_pgd;
    cur_vpn = ((uint64)(va) >> 30) & 0x1fff;
    cur_pte = *(cur_pgtbl + cur_vpn);
    if (!(cur_pte & 1)) {
        return;
    }
    cur_pgtbl = (uint64*)((cur_pte >> 10) << 12) + PA2VA_OFFSET;
    cur_vpn = ((uint64)(va) >> 21) & 0x1fff;
    cur_pte = *(cur_pgtbl + cur_vpn);
    if (!(cur_pte & 1)) {
        return;
    }
    cur_pgtbl = (uint64*)((cur_pte >> 10) << 12) + PA2VA_OFFSET;
    cur_vpn = ((uint64)(va) >> 12) & 0x1fff;
    cur_pte = *(cur_pgtbl + cur_vpn);
    if (!(cur_pte & 1)) {
        return;
    }
    uint64 page = alloc_page();
    create_mapping((uint64)child_pgd, va, (uint64)page-PA2VA_OFFSET, PGSIZE, 31);
    for(int i = 0; i < PGSIZE; i++){
        ((char*)page)[i] = ((char*)va)[i];
    }
    return;
}
```

1.3 修改 `sys_call` 和 `_traps`

`sys_call` 需要额外处理 `SYS_CLONE` 逻辑

```
else if (regs->x[17] == SYS_CLONE) {
    regs->x[10] = sys_clone(regs);
}
```

`_traps` 需要 `__ret_from_fork` 符号, 让child task在做 `schedule` 时能从中断返回后开始执行

```
_traps:
    ...(save registers)
    call trap_handler
__ret_from_fork:
    ...(load registers)
    sret
```

2 编译及测试

给的三个测试点都通过

```
...sup_vm_and done:
[S] Initialized: pid: 1, priority: 1, counter: 0
2022 Hello RISC-V
SET [PID = 1 PRIORITY = 1 COUNTER = 1]
switch to [PID = 1 PRIORITY = 1 COUNTER = 1]
[S] Supervisor Page Fault, scause: 000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010158
[S] New task: 2
[S] Supervisor Page Fault, scause: 000000000000000d, stval: 0000000000011978, sepc: 0
00000000000101f0
[U-PARENT] pid: 1 is running!, global_variable: 0
[U-PARENT] pid: 1 is running!, global_variable: 1
switch to [PID = 2 PRIORITY = 1 COUNTER = 1]
[S] Supervisor Page Fault, scause: 000000000000000d, stval: 0000000000011978, sepc: 0
000000000001018c
[U-CHILD] pid: 2 is running!, global_variable: 0
[U-CHILD] pid: 2 is running!, global_variable: 1
[U-CHILD] pid: 2 is running!, global_variable: 2
```

```

[S] Initialized: pid: 1, priority: 1, counter: 0
2022 Hello RISC-V
SET [PID = 1 PRIORITY = 1 COUNTER = 1]
switch to [PID = 1 PRIORITY = 1 COUNTER = 1]
[S] Supervisor Page Fault, scause: 000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010158
[S] Supervisor Page Fault, scause: 000000000000000d, stval: 0000000000011a00, sepc: 0
000000000001017c
[U] pid: 1 is running!, global_variable: 0
[U] pid: 1 is running!, global_variable: 1
[U] pid: 1 is running!, global_variable: 2
[S] New task: 2
[U-PARENT] pid: 1 is running!, global_variable: 3
[U-PARENT] pid: 1 is running!, global_variable: 4
[U-PARENT] pid: 1 is running!, global_variable: 5
switch to [PID = 2 PRIORITY = 1 COUNTER = 1]
[U-CHILD] pid: 2 is running!, global_variable: 3
[U-CHILD] pid: 2 is running!, global_variable: 4
[U-CHILD] pid: 2 is running!, global_variable: 5

```

```

[S] Initialized: pid: 1, priority: 1, counter: 0
2022 Hello RISC-V
SET [PID = 1 PRIORITY = 1 COUNTER = 1]
switch to [PID = 1 PRIORITY = 1 COUNTER = 1]
[S] Supervisor Page Fault, scause: 000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 000000000000000f, stval: 0000003fffffffff8, sepc: 0
0000000000010158
[S] Supervisor Page Fault, scause: 000000000000000d, stval: 0000000000011930, sepc: 0
0000000000010174
[U] pid: 1 is running!, global_variable: 0
[S] New task: 2
[U] pid: 1 is running!, global_variable: 1
[S] New task: 3
[U] pid: 1 is running!, global_variable: 2
[U] pid: 1 is running!, global_variable: 3
[U] pid: 1 is running!, global_variable: 4
switch to [PID = 2 PRIORITY = 1 COUNTER = 1]
[U] pid: 2 is running!, global_variable: 1
[S] New task: 4
[U] pid: 2 is running!, global_variable: 2
[U] pid: 2 is running!, global_variable: 3
[U] pid: 2 is running!, global_variable: 4
switch to [PID = 3 PRIORITY = 1 COUNTER = 1]
[U] pid: 3 is running!, global_variable: 2
[U] pid: 3 is running!, global_variable: 3
[U] pid: 3 is running!, global_variable: 4
switch to [PID = 4 PRIORITY = 1 COUNTER = 1]

```

斐波那契数列也通过

```

SET [PID = 1 PRIORITY = 1 COUNTER = 1]
switch to [PID = 1 PRIORITY = 1 COUNTER = 1]
[S] Supervisor Page Fault, scause: 0000000000000000c, stval: 00000000000100e8, sepc: 0
00000000000100e8
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000003fffffffff8, sepc: 0
00000000000101e0
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000000000011bd8, sepc: 0
0000000000010214
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000000000012000, sepc: 0
0000000000010214
[S] Supervisor Page Fault, scause: 0000000000000000f, stval: 0000000000013000, sepc: 0
0000000000010214
[S] New task: 2
[U] fork returns 2
[U-PARENT] pid: 1 is running! the 0th fibonacci number is 1 and the number @ 1000 in
the large array is 20
[U-PARENT] pid: 1 is running! the 1th fibonacci number is 1 and the number @ 999 in t
he large array is 999
[U-PARENT] pid: 1 is running! the 2th fibonacci number is 1 and the number @ 998 in t
he large array is 998
[U-PARENT] pid: 1 is running! the 3th fibonacci number is 2 and the number @ 997 in t
he large array is 997
[U-PARENT] pid: 1 is running! the 4th fibonacci number is 3 and the number @ 996 in t
he large array is 996
[U-PARENT] pid: 1 is running! the 5th fibonacci number is 5 and the number @ 995 in t
he large array is 995
[U-PARENT] pid: 1 is running! the 6th fibonacci number is 8 and the number @ 994 in t
he large array is 994

```

```

switch to [PID = 2 PRIORITY = 1 COUNTER = 1]
[U] fork returns 0
[U-CHILD] pid: 2 is running! the 0th fibonacci number is 1 and the number @ 1000 in t
he large array is 20
[U-CHILD] pid: 2 is running! the 1th fibonacci number is 1 and the number @ 999 in th
e large array is 999
[U-CHILD] pid: 2 is running! the 2th fibonacci number is 1 and the number @ 998 in th
e large array is 998
[U-CHILD] pid: 2 is running! the 3th fibonacci number is 2 and the number @ 997 in th
e large array is 997
[U-CHILD] pid: 2 is running! the 4th fibonacci number is 3 and the number @ 996 in th
e large array is 996
[U-CHILD] pid: 2 is running! the 5th fibonacci number is 5 and the number @ 995 in th
e large array is 995
[U-CHILD] pid: 2 is running! the 6th fibonacci number is 8 and the number @ 994 in th
e large array is 994
[U-CHILD] pid: 2 is running! the 7th fibonacci number is 13 and the number @ 993 in t
he large array is 993
[U-CHILD] pid: 2 is running! the 8th fibonacci number is 21 and the number @ 992 in t
he large array is 992
[U-CHILD] pid: 2 is running! the 9th fibonacci number is 34 and the number @ 991 in t
he large array is 991
[U-CHILD] pid: 2 is running! the 10th fibonacci number is 55 and the number @ 990 in
the large array is 990
[U-CHILD] pid: 2 is running! the 11th fibonacci number is 89 and the number @ 989 in

```

3 思考题

1.参考 `task_init` 创建一个新的 `task`，将的 `parent task` 的整个页复制到新创建的 `task_struct` 页上。这一步复制了哪些东西？

复制了将parent task的 `task_struct`：

```
struct task_struct {
    struct thread_info thread_info;
    uint64 state;      // 线程状态
    uint64 counter;    // 运行剩余时间
    uint64 priority;   // 运行优先级 1最低 10最高
    uint64 pid;        // 线程id

    struct thread_struct thread;
    uint64 satp;
    pagetable_t pgd;
    uint64 vma_cnt;    /* 下面这个数组里的元素的数量 */
    struct vm_area_struct vmas[0]; /* 为什么可以开大小为 0 的数组？
                                   这个定义可以和前面的 vma_cnt 换个位置吗？ */
};
```

此时 `pid`, `thread.ra/sp`, `satp`, `pgd` 还是parent task的，同时vma应该映射的部分也还没映射，需要做后续调整。

2.将 `thread.ra` 设置为 `__ret_from_fork`，并正确设置 `thread.sp`。仔细想想，这个应该设置成什么值？可以根据 `child task` 的返回路径来倒推。

应该设置成parent task此时的`sp`，即目前的栈顶：`pt_regs`。因为`child task`的返回路径是：`__switch_to->__ret_from_fork(in_traps)->user program`，因此 `__switch_to` 需要给 `__ret_from_fork` 正确的栈指针，具体还要看每个人 `_traps` 的设计，总之要给 `__ret_from_fork` 执行时正确的栈指针，一般而言是 `pt_regs`。

3.利用参数 `regs` 来计算出 `child task` 的对应的 `pt_regs` 的地址，并将其中的 `a0`, `sp`, `sepc` 设置成正确的值。为什么还要设置 `sp`？

`thread.sp` 是为了保证 `__ret_from_fork` 执行时有正确的栈指针，而设置栈内的`sp`(即`pt_regs->sp`)是为了在 `__ret_from_fork` 的出栈时能够找到正确的`sp`。